

Bahram Behzadian

Palo Alto, CA

Email: behzadian.bahram@gmail.com

Website: buiksat.github.io

LinkedIn: [buiksat](#)

Google Scholar | U.S. permanent resident

Reinforcement Learning | Evaluation | Post-Training

Profile

Research scientist at Meta with a Ph.D. in Computer Science focused on reinforcement learning. I combine research on robust MDPs and policy improvement under imperfect evaluators with production experimentation and evaluation systems. My research on truncated evaluators and verifier-style rewards connects to RL post-training and reasoning.

Experience

2022–Present **Research Scientist, Meta**

- Build production systems for experiment-metric correctness, anomaly and data-quality monitoring, and ML auditing in large-scale ads experimentation.
- Develop statistical diagnostics for unreliable, biased, stale, or inconsistent signals and translate failure analysis into monitoring and correctness safeguards.

2018 **AI Engineer Intern, Envio, Inc., Dover, NH**

- Developed learning-based methods for vehicle routing and scheduling under operational constraints and uncertain execution.

Research Highlights

Policy Improvement under Imperfect Evaluators

UPI-TRM, ICLR 2026 Workshop on AI with Recursive Self-Improvement

Derived finite-depth policy-improvement guarantees under explicit evaluator assumptions, separating Bellman residual from truncation bias; empirically studied policy improvement with verifier-style rewards under truncated evaluators.

Robust and Risk-Sensitive Markov Decision Processes

NeurIPS 2021 and AISTATS 2021

Developed scalable algorithms for L_∞ -constrained S-rectangular robust MDPs and percentile criteria under model uncertainty.

Technical Skills

Programming: Python, PyTorch.

Evaluation and experimentation: experiment-metric correctness, statistical diagnostics, anomaly and data-quality monitoring, ML auditing.

Research: reinforcement learning, robust MDPs, policy improvement, model-based RL, partial observability, truncated evaluators.

Post-training: RL with verifiable rewards (RLVR); working knowledge of PPO, DPO, GRPO, and reward modeling.

Selected Publications

UPI-TRM: Policy Improvement under Truncated Evaluators with Verifier-Style Rewards

Bahram Behzadian, Brett Daley, Gopeshh Subbaraj, Houssam Nassif

ICLR 2026 Workshop on AI with Recursive Self-Improvement [Paper]

Fast Algorithms for L_∞ -Constrained S-Rectangular Robust MDPs

Bahram Behzadian, Marek Petrik, Chin Pang Ho

Conference on Neural Information Processing Systems (NeurIPS), 2021 [Paper]

Optimizing Percentile Criterion using Robust MDPs

Bahram Behzadian, Reazul Hasan Russel, Marek Petrik, Chin Pang Ho

International Conference on Artificial Intelligence and Statistics (AISTATS), 2021 [Paper]

Fast Feature Selection for Linear Value Function Approximation

Bahram Behzadian, Soheil Gharatappeh, Marek Petrik

International Conference on Automated Planning and Scheduling (ICAPS), 2019 [Paper]

Monte Carlo Localization in Hand-Drawn Maps

Bahram Behzadian, Pratik Agarwal, Wolfram Burgard, Gian Diego Tipaldi

IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS), 2015 [Paper]

Complete publication list: [Publications page](#) | [Google Scholar](#)

Education

2015–2022 **Ph.D. in Computer Science**, *University of New Hampshire*

Dissertation: *Efficient Data-Driven Robust Policies for Reinforcement Learning* (advisor: Marek Petrik).

M.S. in Computer Science awarded en route (2018). Thesis: *Representation Learning for Reinforcement Learning via Singular Value Decomposition*.

2010–2013 **M.S. in Machine Automation**, *Tampere University*

Thesis: *Robot Localization under Partial Observability with Weak Maps*

Honors and Fellowships

CEPS Graduate Fellowship, University of New Hampshire

Thesis and Dissertation Fellowship, Tampere University of Technology